

* Architecture: Organization of the hardware components of a computer; Software interface to get the components to carry out desired operation
 * machine language : Binary representation used for communication within a computer system. assembly language: A symbolic representation of machine instructions. GB = 10⁹ GHz

clock = 10⁹

* .text: Subsequent items(instructions) stored in Text segment at next available address.

.globl: Declare the listed label(s) as global to enable referencing from other file.

.data: Subsequent items stored in Data segment at next available address.

.ascii: Store the string in the Data segment and add null terminator.

* MIPS instructions: * op: Basic operation of the instruction, traditionally called the opcode. * rs: The first register source operand.

* rt: The second register source operand. * rd: The register destination operand. It gets the result of the operation.

* shamt: Shift amount. (Section 2.6 explains shift instructions and this term; it will not be used until then, and hence the field contains zero in this section.)

* funct: Function. This field, often called the function code, selects the specific variant of the operation in the op field.

* sys call: part of the application Binary Interface(simplified interface to I/O)

li: load immediate move(move constants of one register to another)/ la: load address

sbrk: unix system called used to implement malloc(allocate space= new in java) and uses jal X to jump to procedure X (sometimes named the callee). The callee then performs the

* call function: the calling program, or caller, puts the parameter values in \$a0-\$a3 and uses jal X to jump to procedure X (sometimes named the callee). The callee then performs the calculations, places the results in \$v0 and \$v1, and returns control to the caller using jr \$ra.

* malloc() allocates space on the heap and returns a pointer to it. free() releases space on the heap to which the pointer points.

* assembly language: A symbolic language that can be translated into binary machine language.

* pseudoinstruction: A common variation of assembly language instructions often treated as if it were an instruction in its own right.

* Given a branch on register \$s0 being equal to register \$s1, beq \$s0, \$s1, L1. Replace it by a pair of instructions that offers a much greater branching distance.

Answer: These instructions replace the short-address conditional branch: bne \$s0, \$s1, L2

j L1

L2:

* Moore's Law. It states that integrated circuit resources double every 18–24 months.

* Amdahl's law: if you speedup part of a system, you are not going to see proportional improvements overall.

* Application software->System software->Hardware

systems software : Software that provides services that are commonly useful, including operating systems, compilers, loaders, and assemblers.

* operating system: Supervising program that manages the resources of a computer for the benefit of the programs that run on that computer.

* compiler: A program that translates high-level language statements into assembly language statements.

* assembler: A program that translates a symbolic version of instructions into the binary version.

Hardware:including I/O, memory,datapath, control.(data path and control combined and called the processor)

* The datapath performs the arithmetic operations, and control tells the datapath, memory, and I/O devices what to do according to the wishes of the instructions of the program.

* memory: The storage area in which programs are kept when they are running and that contains the data needed by the running programs.

* High-level-language(e.g. C)-->(Compiler)-->Assembly language-->(Assembler)-->Binary machine language.

* Instruction set architecture: Also called architecture. An abstract interface between the hardware and the lowest-level software that encompasses all the information necessary to write a machine language program that will run correctly, including instructions, registers, memory access, I/O, and so on.

* Application binary interface (ABI): The user portion of the instruction set plus the operating system interfaces used by application programmers. It defines a standard for binary portability across computers.

* Interpretations of performance: execution time, throughput, power consumption.

* Response time: Also called execution time. The total time required for the computer to complete a task, including disk accesses, memory accesses, I/O activities, operating system overhead, CPU execution time, and so on.

Performance = 1/ Execution time or execution time per program.

1. Wall clock time, response time, or elapsed time: These terms mean the total time to complete a task, including disk accesses, memory accesses, input/output (I/O) activities, operating system overhead—everything.

2. CPU execution time or simply CPU time: The time the CPU spends computing for this task and does not include time spent waiting for I/O or running other programs. The actual time the CPU spends computing for a specific task.

3. Clock cycle: Also called tick, clock tick, clock period, clock, or cycle. The time for one clock period, usually of the processor clock, which runs at a constant rate.

4. Clock period: The length of each clock cycle. Clock rate (e.g., 4 gigahertz, or 4 GHz), which is the inverse of the clock period. = cycles / second

5. CPI(clock cycles per instruction): Average number of clock cycles per instruction for a program or program fragment.= clock cycles/ instruction count

6. IPC(instructions per clock cycle): Average number of instructions per clock cycle for a program or program fragment.

7. CPU time = Instruction Count * CPI * Clock cycle time = Instruction Count * CPI / Clock rate

CPU time = Clock Cycles * Clock Cycle time = Clock cycles / Clock rate

* Throughput: Also called bandwidth. Another measure of performance, it is the number of tasks completed per unit time.

* Power = Energy * Capacitive load * Voltage²

Components affecting performance: 1. algorithm: Instruction count, (CPI)2. programming language: Instruction count, CPI 3. compiler: Instruction count, CPI 4.processor 5. I/O 6.Instruction set architecture: Instruction count, CPI, clock rate

* SPEC (System Performance Evaluation Cooperative) is an effort funded and supported by a number of computer vendors to create standard sets of benchmarks for modern computer systems.

Dividing the execution time of a reference processor by the execution time of the measured computer normalizes the execution time measurements; this normalization yields a measure, called the SPECratio, which has the advantage that bigger numeric results indicate faster performance.

* MIPS (million instructions per second): A measurement of program execution speed based on the number of millions of instructions. MIPS is computed as the instruction count divided by the product of the execution time and 10⁶ = Instruction count / (Execution time*10⁶) = Clock rate / (CPI*10⁶)

* MFLOPS: millions of floating point ops per sec

Deci	bin	hex	oct
0	0000	0	00
1	0001	1	01
2	0010	2	02
3	0011	3	03
4	0100	4	04
5	0101	5	05
6	0110	6	06
7	0111	7	07
8	1000	8	10
9	1001	9	11
10	1010	a	12
11	1011	b	13
12	1100	c	14
13	1101	d	15
14	1110	e	16
15	1111	f	17

Purpose: Implement recursive binary search in MIPS32

Input: n: number of elements in list; list: sorted list of n ints; val: int to search for

Output: Subscript of val if val is in list ; Otherwise -1

.text
 .globl main

main:

addi \$sp, \$sp, -416 # Make additional stack space. 4 words for \$ra, \$s0, \$s1, \$s2.100 words for list

sw \$ra, 412(\$sp) # Put contents of \$ra on stack

sw \$s0, 408(\$sp) # Put \$s0 on stack

sw \$s1, 404(\$sp) # Put \$s1 on stack

sw \$s2, 400(\$sp) # Put \$s2 on stack

move \$s0, \$sp # \$s0 = stores start address of list = \$sp

Ask the OS to read a number and put it in \$s1 = n

li \$v0, 5 # Code for read int.

syscall # Ask the system for service.

move \$s1, \$v0 # Put the input value (n) in a safe place

Now read in the list

move \$a0, \$s0 # First arg is list

move \$a1, \$s1 # Second arg is n

jal rd_list

Read the value to search for, val

li \$v0, 5 # Code for read int.

syscall # Ask the system for service.

move \$s2, \$v0 # Put the input value (val) in a safe place

Now carry out the binary search

move \$a0, \$s2 # First arg is val

move \$a1, \$s0 # Second arg is list

move \$a2, \$zero # Third arg is 0

```

    addi $a3, $s1, -1      # Fourth arg is n-1
    jal  binsrch
    # Now print the result
    move $a0, $v0          # Arg is return val from bin_srch
    li   $v0, 1            # Code for print int
    syscall
    # Prepare for return
    lw   $ra, 412($sp)      # Retrieve return address
    lw   $s0, 408($sp)      # Retrieve $s0
    lw   $s1, 404($sp)      # Retrieve $s1
    lw   $s2, 400($sp)      # Retrieve $s2
    addi $sp, $sp, 416      # Make additional stack space.
    li   $v0, 10           # For MARS
    syscall
    ### Read_list function
    # $a0 is the address of the beginning of list (In/out); $a1 is n (In)
rd_list:
    # Setup
    addi $sp, $sp, -4      # Make space for return address
    sw   $ra, 0($sp)       # Save return address
    # Main for loop
    move $t0, $zero        # $t0 = i = 0
rd_list: bge $t0, $a1, rddone # If i = $t0 >= $a1 = n branch out of loop. Otherwise continue.
    li   $v0, 5            # Code for read int.
    syscall                # Ask the system for service.
    sll  $t1, $t0, 2        # Words are 4 bytes: use 4*i, not i
    add  $t1, $a0, $t1      # $t1 = list + i
    sw   $v0, 0($t1)       # Put the input value in $v0 in list[i]
    addi $t0, $t0, 1        # i++
    j    rd_list           # Go to the loop test
    # Prepare for return
rddone: lw   $ra, 0($sp)    # retrieve return address
    addi $sp, $sp, 4        # adjust stack pointer
    jr   $ra               # return
    ### Print_list function (Only for Debugging): $a0 is the address of the beginning of list (In), $a1 is n (In).
pr_list:
    # Setup
    addi $sp, $sp, -4      # Make space for return address
    sw   $ra, 0($sp)       # Save return address
    # Main for loop
    move $t2, $a0          # Need $a0 for syscall: so copy to t2
    move $t0, $zero        # $t0 = i = 0
pr_list: bge $t0, $a1, prdone # If i = $t0 >= $a1 = n branch out of loop. Otherwise continue.
    sll  $t1, $t0, 2        # Words are 4 bytes: use 4*i, not i
    add  $t1, $t2, $t1      # $t1 = list + i
    lw   $a0, 0($t1)       # Put the value to print in $a0
    li   $v0, 1            # Code for print int.
    syscall
    # Print a space
    la   $a0, space        #
    li   $v0, 4            # Code for print string
    syscall
    addi $t0, $t0, 1        # i++
    j    pr_list           # Go to the loop test
    # print a newline
prdone: la   $a0, newln
    li   $v0, 4            # code for print string
    syscall
    # Prepare for return
    lw   $ra, 0($sp)       # retrieve return address
    addi $sp, $sp, 4        # adjust stack pointer
    jr   $ra               # return
    ### Bin_srch function
    # $a0 is the value to be searched for (val), $a1 is the list, $a2 is the current lower bound (lo), $a3 is the current upper bound (hi)
    # All args are input args: $a0 and $a1 are unchanged in the recursive calls. $a2 and $a3 are changed, but after a call starts a recursive call, their current values aren't used again.
binsrch: # Bin_srch function
    # Setup
    addi $sp, $sp, -4      # Make space for return address
    sw   $ra, 0($sp)       # Save return address
    bgt  $a2, $a3, not_in  # If lo = $a2 > $a3 = hi, val isn't in the list, quit
    add  $t0, $a2, $a3      # $t0 = hi + lo
    srl  $t0, $t0, 1        # $t0 = mid = $t0/2
    sll  $t1, $t0, 2        # $t1 = byte offset of list[mid]
    add  $t1, $t1, $a1      # $t1 = absolute addr of list[mid]
    lw   $t2, 0($t1)       # $t2 = list[mid]
    bne  $a0, $t2, not_eq   # If $a0 = val != list[mid] = $t2, go to not_eq
    # We found val
    move $v0, $t0          # val = list[mid], set return value to mid.
    j    done
    # We didn't find val this time
not_eq: bgt  $a0, $t2, grter # if val > list[mid] go to grter
    # val < list[mid], update hi
less:  addi $a3, $t0, -1    # hi = mid-1
    jal  binsrch
    j    done              # Don't forget this!
    # val > list[mid], update lo
grter: addi $a2, $t0, 1     # lo = mid+1
    jal  binsrch
    j    done              # Don't forget this!
    # val is not in list
not_in: li   $v0, -1       # val not in list: return -1
    # Prepare for return
done:  lw   $ra, 0($sp)    # Retrieve return address
    addi $sp, $sp, 4        # Adjust stack pointer
    jr   $ra               # return
.data
space: .asciiz " "
newln: .asciiz "\n"

```