



Artificial Intelligence Programming

Information Retrieval

Cindi Thompson

Department of Computer Science
University of San Francisco

Information Retrieval

Earlier we said that approaches to working with natural language can be divided roughly into two camps:

- Natural Language Processing (or Analysis)
- Information Retrieval
 - Focus on statistical summarization of a document.
 - Shallower analysis, much more scalable.

Today we'll cover the latter (IR)

Information Retrieval

- Information retrieval deals with the storage, retrieval, organization of, and access to information items
- Overlaps with:
 - Databases (more of a focus on content)
 - AI
 - Search engines

Needs and queries

- A user typically has an *information need*.
- The job of an IR system is to translate that need into a query language and then find documents that satisfy that need.
- What are some sorts of query languages?

Query Languages

- What are some sorts of query languages?
- Keyword - Google, Yahoo!, etc.
- Natural language - Ask.com
- SQL-style
- Similar item - Netflix, Amazon
- Multimedia - Flickr

User tasks

- We'll also distinguish between different types of user tasks.
- The most common are *searching* and *browsing*.
 - Searching - the user has a specific information need, and wants a document that meets that need.
 - “Find me an explanation of the re module in Python”
 - Browsing - the user has a broadly defined set of interests, and wants information that satisfies his/her interests.
 - “Find me interesting pages about Python”
- These different modes have different models of success.

User tasks

- Searching and browsing are both *pull* tasks.
 - User is actively fetching information from a repository.
- We can also think about *push* tasks, where selected data is delivered to a client as it is made available.
- This is called *filtering*
 - RSS readers are an example of this, as is Google News.

Modeling a Document

- In order to match a query to a document, an IR system must have a *model* of the document.
- This might be:
 - A category or description (as in a library)
 - A set of extracted phrases or keywords
 - The full text of the document
 - Full text with filtering

“Bag of words” model

- The techniques we’ll look at today treat a document as a *bag of words*.
- Order is discarded; we just count how often each word appears.
- No semantics involved
- Intuition: Frequently-appearing words give an indication of subject matter.
- Advantage: No need to parse, computationally tractable for large collections.
- Disadvantage: Contextual information and meaning is lost.

Data cleaning

- When preparing a document such as a webpage for an IR system, the data must typically be *cleaned* first.
 - HTML, Javascript removed.
 - (Links and structural information might be kept separately)
 - Non-words removed.
 - Converted to lower case
 - *stopwords* removed. These are words that have little or no semantic content. (a, an, the, he, she, their, among, etc)

Data Cleaning

- *Stemming* might also be performed.
- Word suffixes, such as pluralization, past tense, -ing are removed.
- run, runs, running, runner all become run.
- Advantages: If we're just counting words, this lets us correctly count different forms of a word.
- Disadvantages: dealing with abnormal forms (person/people, run/ran), potential misgrouping (university, universal)
- The stemmer can be tuned to minimize either *false positives* (accidentally stemming a word it shouldn't) or *false negatives* (not stemming a word it should.)
- There's some debate in the research community about the effectiveness of stemming.

“Bag of Words”

- Once a document has been cleaned, the simplest model just counts how many times each word occurs in the document.
- This is typically represented as a dictionary.
- You built this in assignment 1.

Evaluating an IR System

- Our prototypical use case is this:
 - The user submits a query to the system
 - Some documents are returned
- How can we evaluate performance?

Precision and Recall

- Precision measures how well returned documents match a query.
 - $\text{precision} = \frac{\text{matchingDocs}}{\text{totalDocsReturned}}$
- Recall measures the fraction of relevant documents returned.
 - $\text{recall} = \frac{\text{relevantDocsReturned}}{\text{totalRelevantDocs}}$
- When might we want high precision? High recall?
- Often, we can trade precision against recall.

Precision and Accuracy more generally

- We can apply these ideas more generally in machine learning as well. In learning, precision refers to the avoidance of false positives.
 - If our learning algorithm says an instance is of class x , how likely is it that that instance truly *is* of class x .
- Recall is the avoidance of false negatives.
 - If our algorithm says an instance is not of class x , how likely is this to be the case?
- In IR, class x is “documents matching our query.”

Boolean Queries

- Boolean queries are simple, but not very practical.
- User provides a set of keywords.
 - Possibly also OR terms
- All documents containing all keywords are returned.
- This is the sort of query model that databases use

Boolean Queries

- Weaknesses:
 - No concept of partial match, or ability to rank results
 - Low recall
 - Boolean queries are awkward for users

Probabilistic Queries

- A simple extension is to allow partial matches on queries
- Score documents according to the fraction of query terms matched
- Return documents according to score
 - Example: Document contains “cat cat dog bunny fish”
 - Query is “cat dog (bunny OR snake) bird”
 - Score is 3/4.

Probabilistic Queries

- Weaknesses:
 - Still requires logical queries
 - Doesn't deal with word frequency
 - Dependent on query length - short queries will have a hard time getting differentiated scores.
 - The average Google query is only three words long!

Dealing with Word Frequency

- Intuitively, some words in a document should matter more than others.
 - The word “aardvark” occurring 10 times in a document is probably more meaningful than the word “date” occurring 10 times.
- We want to weight words such that words which are rare in general, but common in a document, are more highly considered.

Building a corpus

- To measure how frequently words occur in general, we must construct a corpus.
 - This is a large collection of documents
- Must be careful to ensure that we select documents of the appropriate style
- Different types of documents have different word frequencies
 - New York Times vs Livejournal
- Recall from NLP discussion: The statistical distribution of words in a corpus is a type of *language model*.

Building a corpus

- We begin by cleaning the data as before
- Construct a dictionary that maps words to the number of pages they occur in.
 - Don't worry about multiple occurrences within a document
- The result is referred to as *document frequency*

TFIDF

- We can now weight each word to indicate its importance in the language model.
- The most common weighting scheme is TF-IDF: term frequency - inverse document frequency.
- $TFIDF(word) = TF(word) * \log(\frac{|corpus|}{DF(word)})$
- $TF(word)$ is how frequently the word occurs in the search query (or a specific document)
- $DF(word)$ is the number of pages in the corpus that contain the word.

TFIDF

- Think about extrema:
 - What happens if a word occurs in exactly one document in the corpus?
 - What happens if a word occurs in every document in the corpus?
- We want to favor words that discriminate interesting pages from non-interesting pages

Word Weighting

- We can now process each document and assign a weight to each word.
- We could use this to improve the performance of the probabilistic scorer.
- More interestingly, we can use it to determine how similar two documents are.
- This gives us another way for users to search
 - “Find more documents like this”

Documents as vectors

- At this point, each document can be represented as a dictionary of words and TFIDF scores

`cat: 4.33; dog: 2.1 ; bunny: 8.2; fish: 0.33`

- Conceptually, these documents can be thought of as an n -dimensional vector, where n is the number of words in the lexicon (all words in all documents) and the value of $v[n]$ is the TFIDF score for that word.
- Many elements of the vector are zero, since those words don't appear in that specific document.

Comparing vectors

- We can now use well-known techniques from geometry to compare these vectors.
- We could measure the angle between the vectors.
 - The scale is not convenient, and the calculation is complicated.
- Easier is to measure the cosine of this angle.
- Identical documents have a cosine of 1, and completely dissimilar documents have a cosine of zero.

Computing cosine similarity

- The formula for the cosine of the angle between two vectors is: $\frac{a \cdot b}{||a|| ||b||}$
 - This is the dot product of the two vectors, divided by the product of their magnitudes.
 - The dot product is computed by summing the product of the respective elements of each vector:
- $$\sum_i v1[i] * v2[i]$$
- The magnitudes are computed by calculating the square root of the sum of the squares of each component. (this is Pythagoras' rule)

- $\sqrt{\sum_i v[i]^2}$

Computing cosine similarity

- The entire formula, in terms of words in documents, looks like this:

$$\cos(d_1, d_2) = \frac{\sum_{word \in d_1 \cap d_2} d_1[word] * d_2[word]}{\sqrt{\sum_{word \in d_1} d_1[word]^2} * \sqrt{\sum_{word \in d_2} d_2[word]^2}}$$

- This is a very powerful and useful technique for comparing documents.
- It can also be used to compare a query to a document.
- We'll return to it when we study clustering.

Putting it together

- To use a vector model:
 - Collect and clean a corpus and compute document frequencies.
 - For each document in the collection, clean and compute document frequencies.
 - For a query or sample document, compute TFIDF scores.
 - Compute cosine similarity for each document in the collection and return results from highest to lowest.

Querying

- So how does this work with querying?
 - User provides one or more documents she likes. We'll call this the query set.
 - Form a query vector out of the query set.
 - System compares the query vector to documents in the query set and returns matches.
 - Top N, or all within a threshold, or all in a category.

Strengths and Weaknesses

● Advantages:

- No need for users to understand a query language. They just need to know “like” and “dislike”
- Can take advantage of frequency of terms - higher information terms get more weight.

● Disadvantages:

- Requires extra preprocessing of data to be searched.
- Users must label documents they like
- Users need to be careful to only label similar documents as “liked.”

Summary

- Searching vs browsing
- “bag of words” model
- Precision and Recall
- Boolean and Probabilistic Queries
- Term Weighting
- Vector Models and cosine similarity