

# Artificial Intelligence Programming

## *Adversarial Search*

Cindi Thompson

Department of Computer Science  
University of San Francisco

# AC-3 Pseudocode

`v` = the variables in our problem.

`d[v]` is the list of values in the domain of each `v`

for vertex in `v` :

    neighbors = all vertices in `v` that share a constraint with vertex

    for `n` in neighbors :

        for value in `d[vertex]` :

            if there is no value in `d[n]` consistent with value:

                remove value from `d[vertex]`

            if `d[vertex]` is empty, return failure

repeat until `d[vertex]` does not change for any `v`

# Overview

- Example games (board splitting, chess, Othello)
- Min/Max trees
- Alpha-Beta Pruning
- Evaluation Functions
- Stopping the Search
- Playing with chance

# Games as Search

"Unpredictable" opponent → specify a move for every possible opponent reply

Time limits → unlikely to find goal, must approximate

Let's start with deterministic, 2-player games

# Two player games

- Board-Splitting Game
  - Two players,  $V$  &  $H$
  - $V$  splits the board vertically, selects one half
  - $H$  splits the board horizontally, selects one half
  - $V$  tries to maximize the final value,  $H$  tries to minimize the final value

14	5	11	4
12	13	9	7
15	13	10	8
16	1	6	2

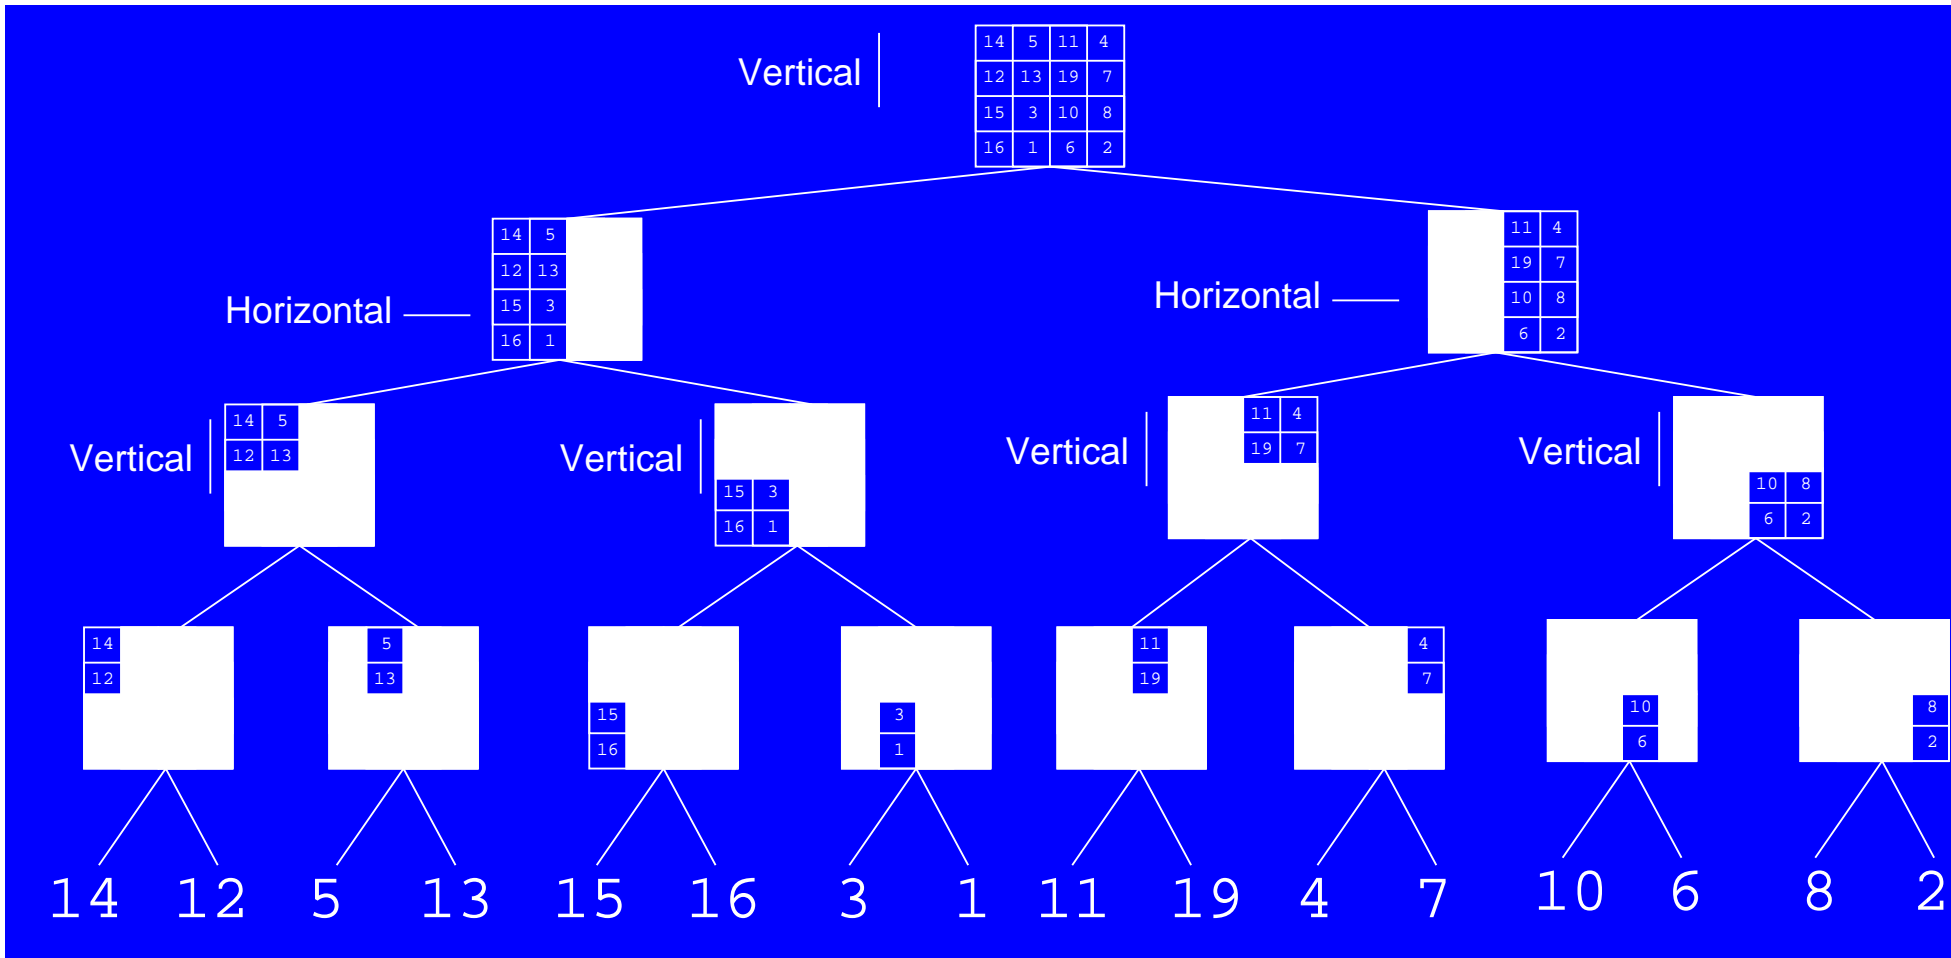
# Two player games

- Board-Splitting Game
  - We assume that both players are rational (make the best possible move)
  - How can we determine who will win the game?
  - And, how can we determine the best move at each state?

# Two player games

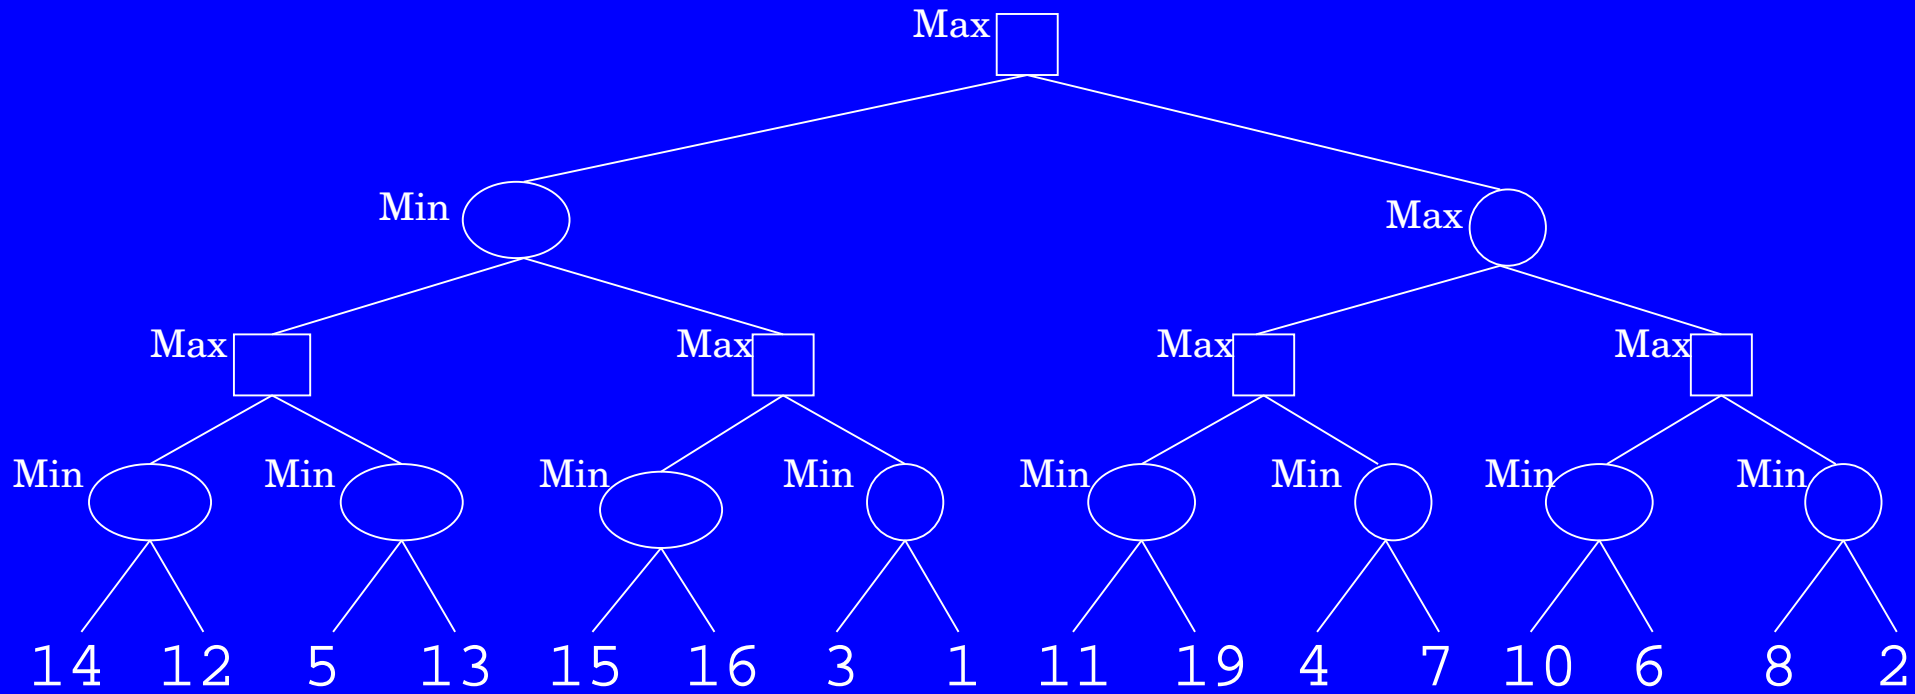
- Board-Splitting Game
  - We assume that both players are rational (make the best possible move)
  - How can we determine who will win the game?
    - Examine all possible games!

# Two player games

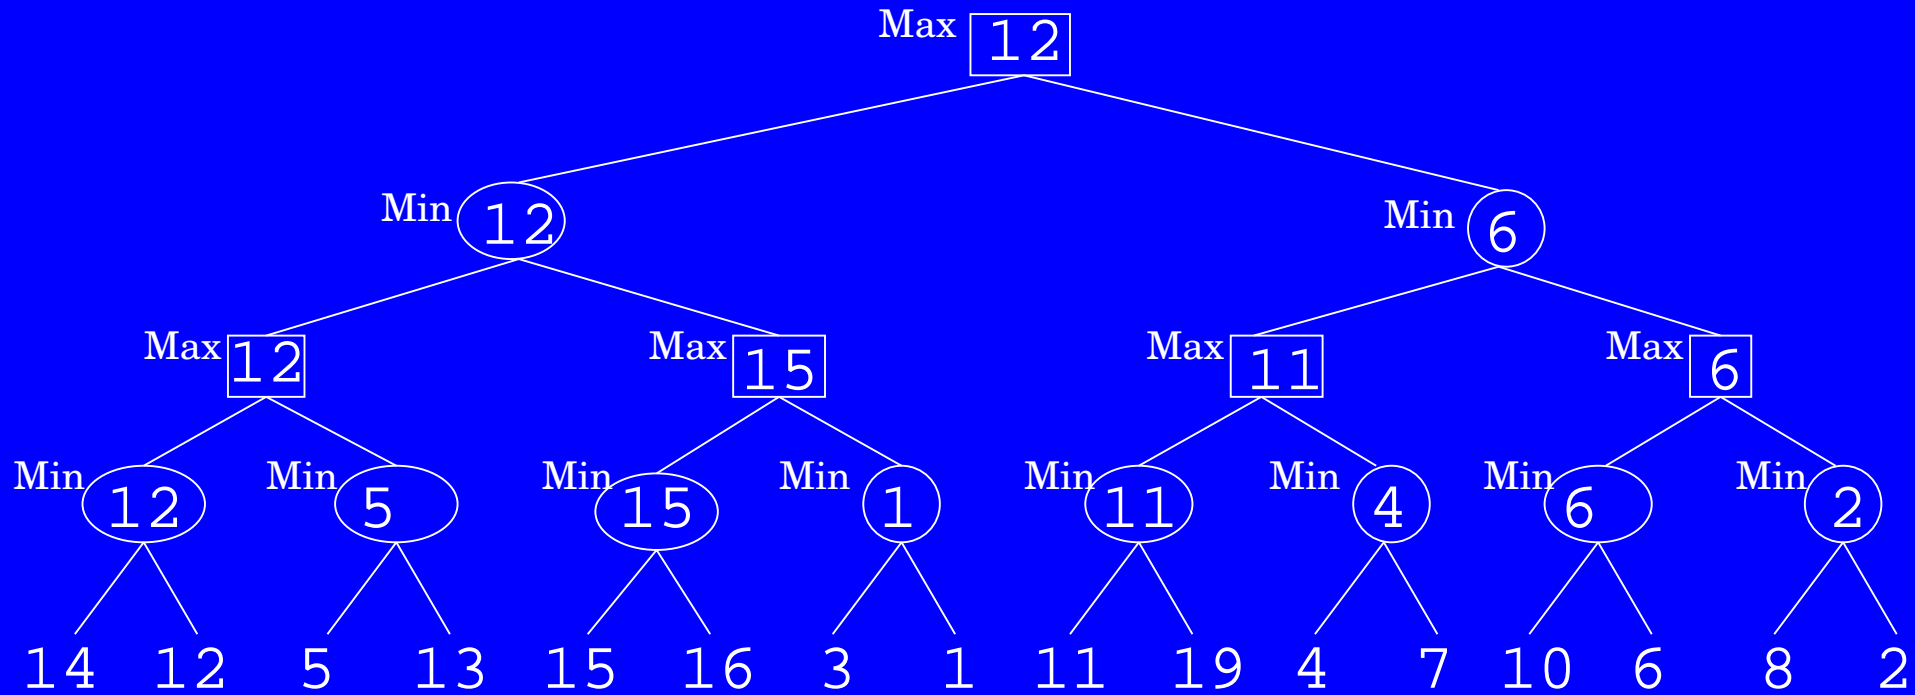




# Two player games



# Two player games



# Two player games

- Game playing agent can do this to figure out which move to make
  - Examine all possible moves
  - Examine all possible responses to each move
  - ... all the way to the last move
  - Calculate the value of each move (assuming opponent plays perfectly)

# Two-Player Games

- Initial state
- Successor Function
  - Just like other Searches
- Terminal Test
  - When is the game over?
- Utility Function
  - Only applies to terminal states
  - Chess: +1, 0, -1
  - Backgammon: 192 ... -192

# Minimax Algorithm

```
def Max-val(node):  
    if terminal(node):  
        return utility(node)  
    maxVal = MIN_VALUE  
    children = successors(node)  
    for child in children:  
        maxVal = max(maxVal, Min-val(child))  
    return maxVal
```

```
def: Min-val(node)  
    if terminal(node):  
        return utility(node)  
    minVal = MAX_VALUE  
    children = successors(node)  
    for child in children:  
        minVal = min(minVal, Max-val(child))  
    return minVal
```

# > 2 Player Games

- What if there are > 2 players?
- We can use the same search tree:
  - Alternate between several players
  - Need a different evaluation function

# > 2 Player Games

- Functions return a vector of utilities
  - One value for each player
  - Each player tries to maximize their utility
  - May or may not be zero-sum

# > 2 Player Games

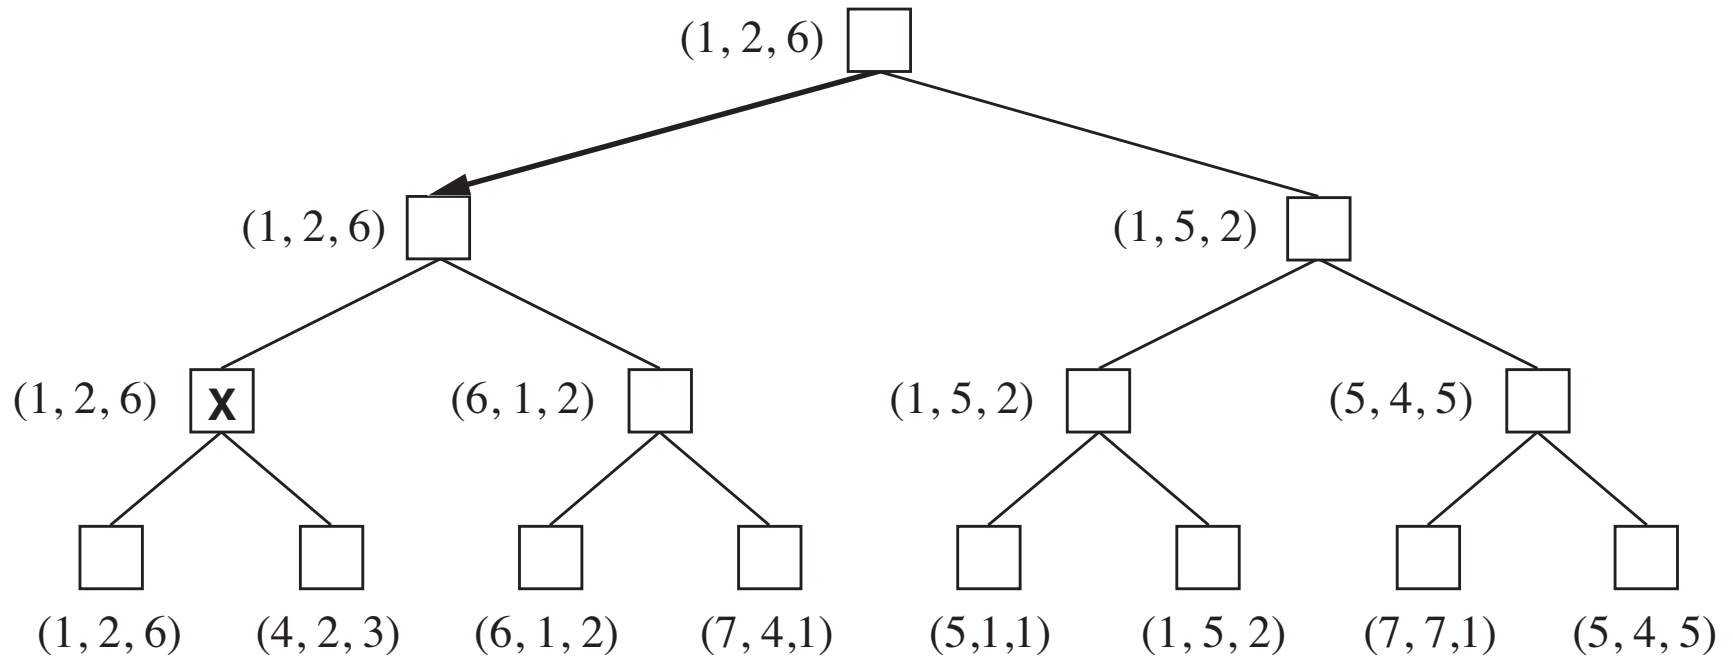
to move

A

B

C

A





# Non zero-sum games

- Even 2-player games don't need to be zero-sum
  - Utility function returns a vector
  - Each player tries to maximize their utility
- If there is a state with maximal outcome for both players, rational players will cooperate to find it
- Minimax is rational, will find such a state

# Minimax Algorithm

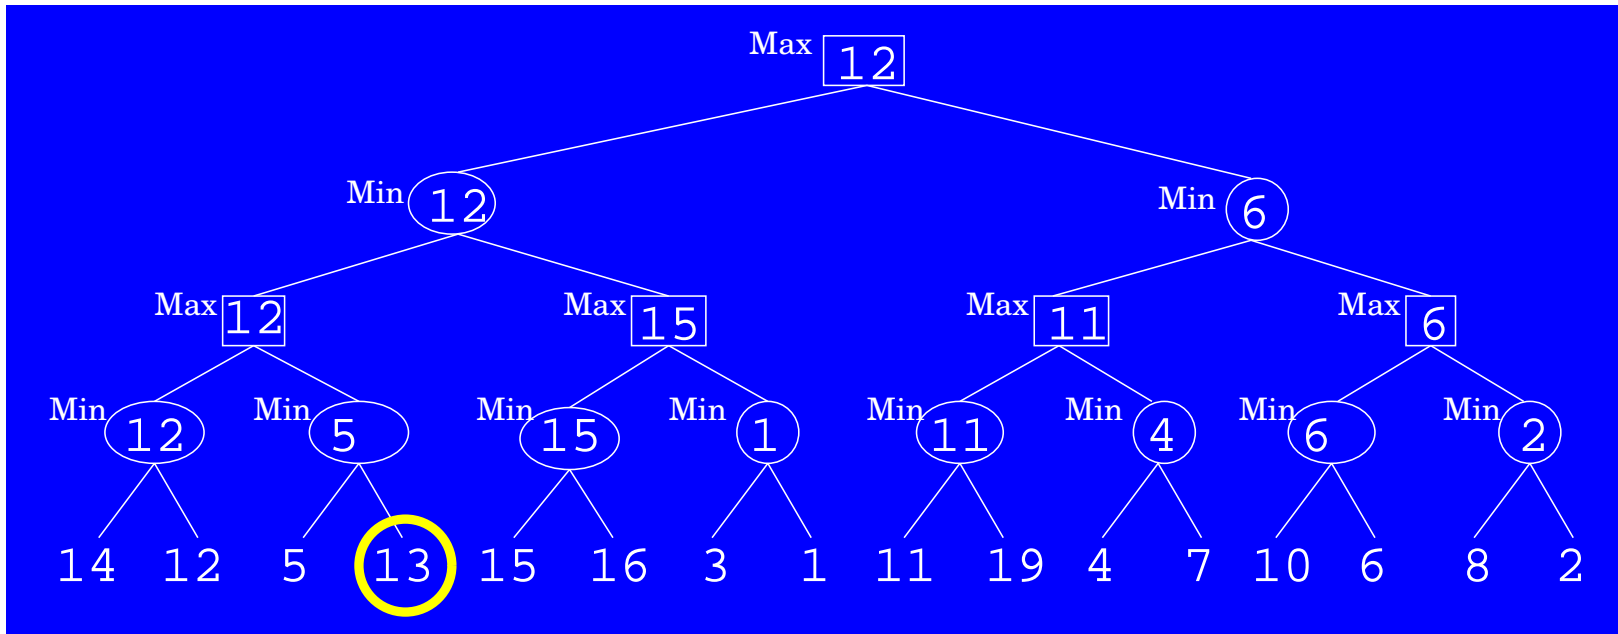
- Complete?
- Optimal?
- Branching factor of  $b$ , game length of  $d$  moves, what are the time and space requirements for Minimax?

# Minimax Algorithm

- Complete? Yes, if tree is finite
- Optimal? Yes, against an optimal opponent
- Branching factor of  $b$ , game length of  $d$  moves, what are the time and space requirements for Minimax?
  - Time:  $O(b^d)$
  - Space:  $O(d)$
- Not manageable for any real games – chess has an average  $b$  of 35, can't search the entire tree
- Need to make this more manageable

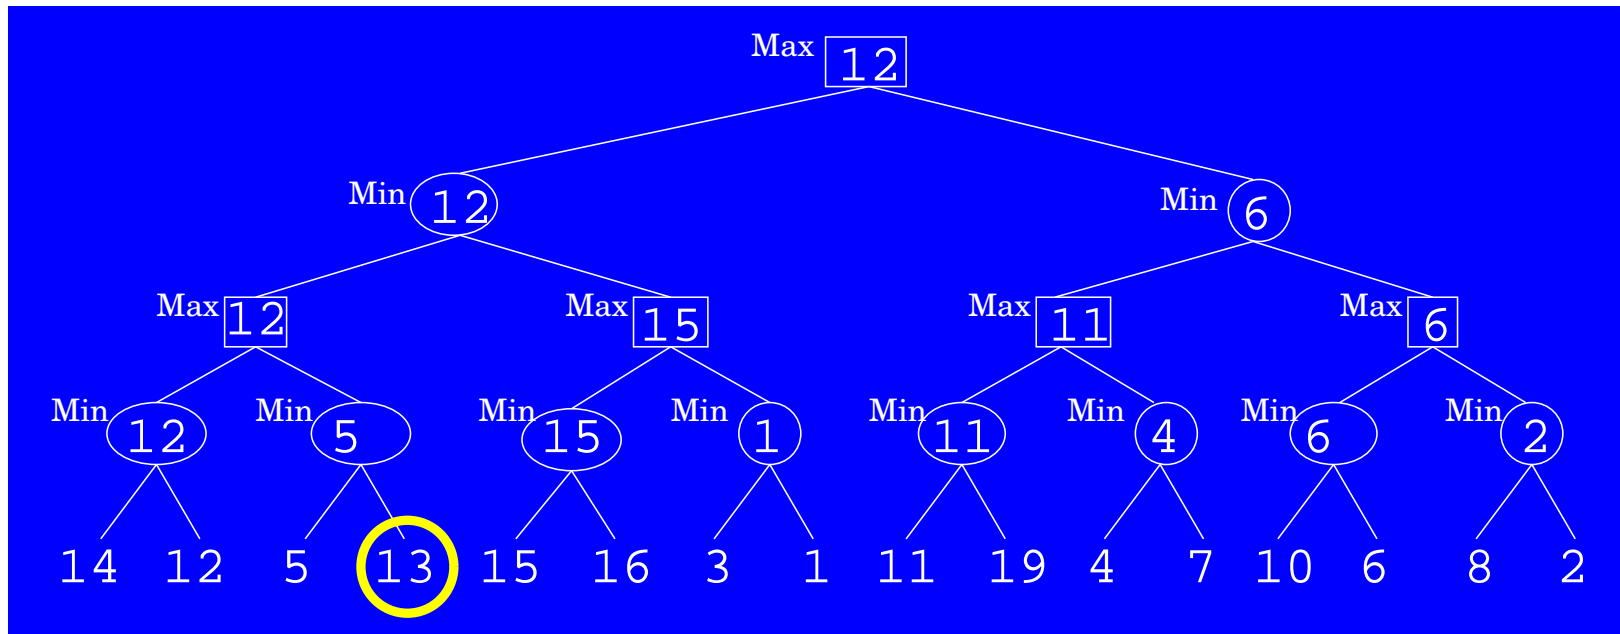
# Alpha-Beta Pruning

- Does it matter what value is in the yellow circle?




# Alpha-Beta Pruning

- If the yellow leaf has a value  $> 5$ , parent won't pick it
- If the yellow leaf has a value  $< 12$ , grandparent won't pick it
- To affect the root, value must be  $< 5$  **and**  $> 12$




# Alpha-Beta Pruning

- Value of nodes in neither yellow circle matter. Are there more?




# Alpha-Beta Pruning

- Value of nodes in none of the yellow circles matter.



# Alpha-Beta Pruning







# Alpha-Beta Pruning

- Maintain two bounds, lower bound  $\alpha$ , and an upper bound  $\beta$ 
  - Bounds represent the values the node must have to possibly affect the root
- As you search the tree, update the bounds
  - Max nodes increase  $\alpha$ , min nodes decrease  $\beta$
- If the bounds ever cross, this branch cannot affect the root, we can prune it.


# Alpha-Beta Pruning




# Alpha-Beta Pruning




# Alpha-Beta Pruning




# Alpha-Beta Pruning




# Alpha-Beta Pruning



# Alpha-Beta Pruning




# Alpha-Beta Pruning







# Alpha-Beta Pruning




# Alpha-Beta Pruning




# Alpha-Beta Pruning




# Alpha-Beta Pruning




# Alpha-Beta Pruning




# Alpha-Beta Pruning



# Alpha-Beta Pruning



# Alpha-Beta Pruning





# Alpha-Beta Pruning

- We can cut large branches from the search tree *without* affecting final result
  - In the previous example, what would happen with similar values and a deeper tree?
- If we choose the order that we evaluate nodes (more on this in a minute...), we can dramatically cut down on how much we need to search

# Evaluation Functions

- We can't search all the way to the bottom of the search tree
  - Trees are just too big
- Search a few levels down, use an evaluation function to see how good the board looks at the moment
- Back up the result of the evaluation function, as if it was the utility function for the end of the game

# Evaluation Functions

- Chess:
  - Material - value for each piece (pawn = 1, bishop = 3, etc)
    - Sum of my material - sum of your material
  - Positional advantages
    - King protected
    - Pawn structure
- Othello:
  - Material – each piece has unit value
  - Positional advantages
    - Edges are good
    - Corners are better
    - “near” edges are bad

# Evaluation Functions

- If we have an evaluation function that tells us how good a move is, why do we need to search at all?
  - Could just use the evaluation function
- If we are only using the evaluation function, does search do us any good?

# Evaluation Functions & $\alpha$ - $\beta$

- How can we use the evaluation function to maximize the pruning in alpha-beta pruning?

# Evaluation Functions & $\alpha$ - $\beta$

- How can we use the evaluation function to maximize the pruning in alpha-beta pruning?
  - Order children of max nodes, from highest to lowest
  - Order children of min node, from lowest to highest
  - (Other than for ordering, eval function is not used for interior nodes)
- With perfect ordering, we need to search only  $b^{d/2}$  (instead of  $b^d$ ) to find the optimal move – can search up to twice as far

A simple example of the value of reasoning about which computations are relevant (a form of *meta-reasoning*)

# Stopping the Search

Still exponential!


- We can't search all the way to the endgame
  - Not enough time
- Search a set number of moves ahead
  - Problems?

# Stopping the Search


- We can't search all the way to the endgame
  - Not enough time
- Search a set number of moves ahead
  - What if we are in the middle of a piece trade?
  - In general, what if our opponent is about to capture one of our pieces



# Stopping the Search



(a) White to move



(b) White to move


# Stopping the Search

- Quiescence Search
  - Only apply the evaluation function to nodes that do not swing wildly in value
  - If the next move makes a large change to the evaluation function, look ahead a few more moves
  - Not increasing the search depth for the entire tree, just around where the action is
  - To prevent the search from going too deep, may restrict the kinds of moves (captures only, for instance)

# Stopping the Search

- Horizon Problem
  - Sometimes, we can push a bad move past the horizon of our search
  - Not preventing the bad move, just delaying it
  - A position will look good, even though it is ultimately bad

# Horizon Problem



# Horizon Problem

- Singular Extensions
  - When we are going to stop, see if there is one move that is clearly better than all of the others.
  - If so, do a quick “search”, looking only at the best move for each player
  - Stop when there is no “clearly better” move
  - Helps with the horizon problem, for a series of forced moves
- Similar to quiescence search

# Cutting off search

- Minimal change to alpha-beta search
- Does it work in practice?
  - $B^m = 10^6, b = 35 \rightarrow m = 4$
- 4-ply look ahead is a hopeless chess player
  - 4-ply  $\approx$  human novice
  - 8-ply  $\approx$  typical PC, human master
  - 12-ply  $\approx$  Deep Blue, Kasparov

# State of the art

- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a precomputed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.

# State of the art

- Othello: human champions refuse to compete against computers, who are too *good*.
- Go: human champions refuse to compete against computers, who are too *bad*. In go,  $b > 300$ , so most programs use pattern knowledge bases to suggest plausible moves.



# Adding Chance

- What about games that have an element of chance (backgammon, poker, etc)
- We can add chance nodes to our search tree
  - Consider “chance” to be another player
- How should we back up values from chance nodes?

# Adding Chance

MAX

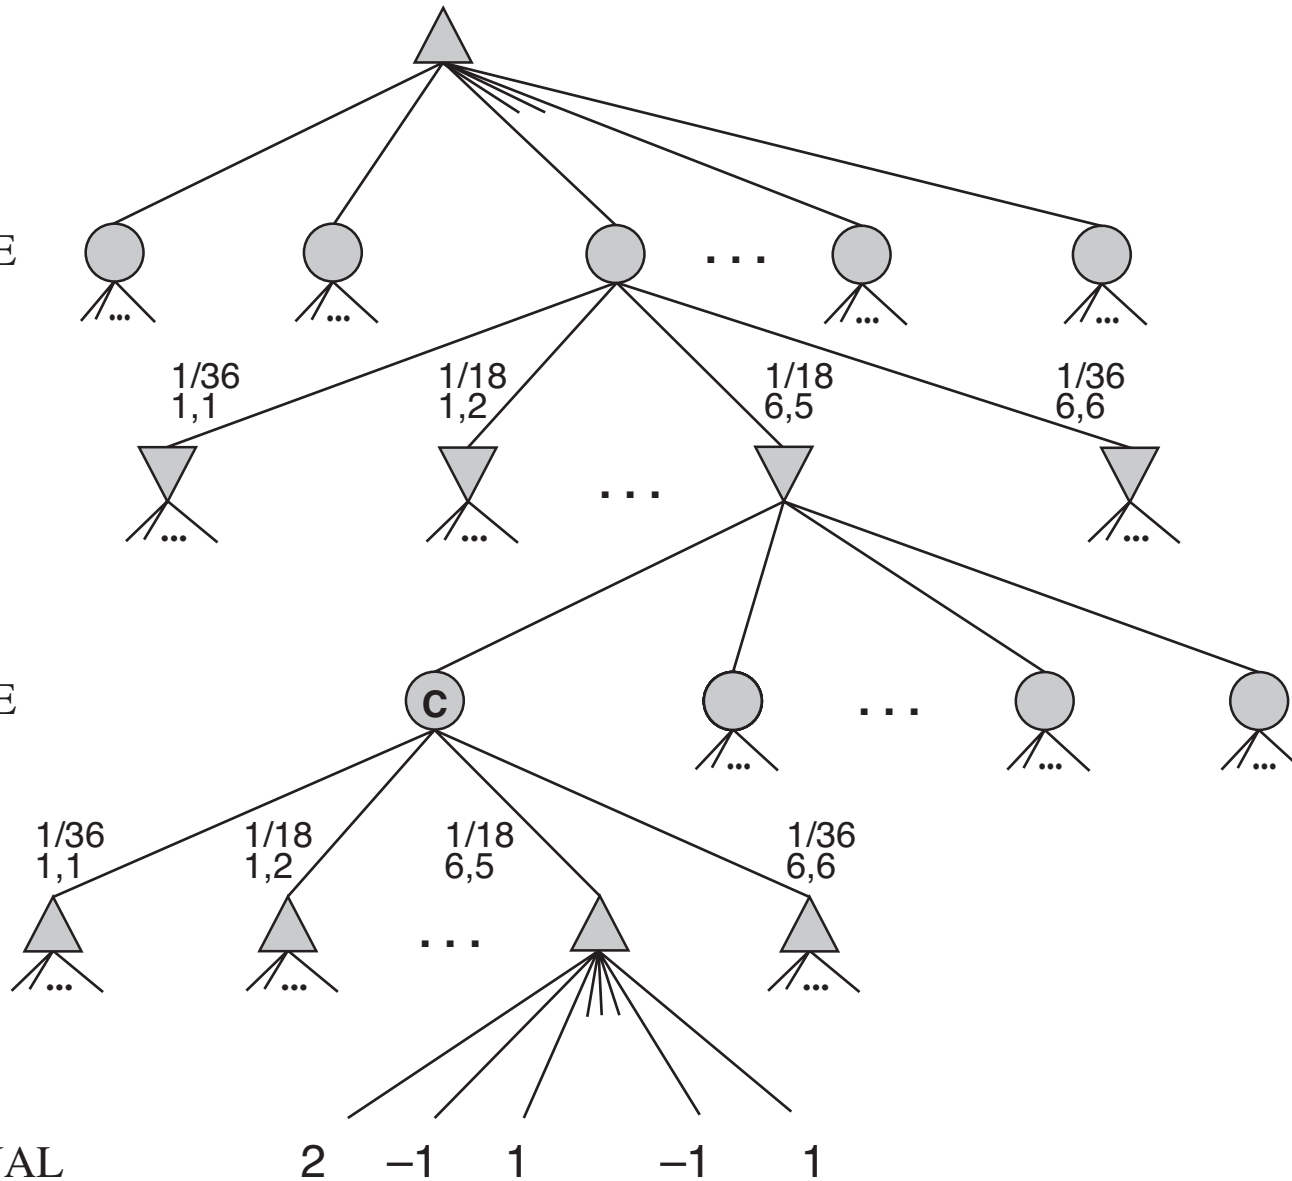
CHANCE

MIN

CHANCE

MAX

TERMINAL



# Adding Chance

- For Max nodes, we backed up the largest value:

$$\max_{s \in \text{Successors}(n)} Val(s)$$

- For Min nodes, we backed up the smallest

$$\min_{s \in \text{Successors}(n)} Val(s)$$

- For chance nodes, we back up the expected value of the node

$$\sum_{s \in \text{Successors}(n)} P(s) Val(s)$$

# Adding Chance

- Adding chance dramatically increases the number of nodes to search
  - Branching factor  $b$  (ignoring die rolls)
  - $n$  different dice outcomes per turn
  - Time to search to level  $m$ ?

# Adding Chance

- Adding chance dramatically increases the number of nodes to search
  - Branching factor  $b$  (ignoring die rolls)
  - $n$  different dice outcomes per turn
  - Time to search to level  $m$ :  $b^m n^m$

# Summary

- Games are fun to work on!
- They illustrate several important points about AI
  - perfection is unattainable: must approximate
  - good idea to think about what to think about

# Summary

- Min/Max trees
- Alpha-Beta Pruning
- Evaluation Functions
- Stopping the Search
- Playing with chance

# Alpha-beta pseudocode, pt 1

```
def alpha-beta-search(state):  
    v = max-val(state, -INF, INF)  
    return action associated with v  
  
def max-val(s, alpha, beta): #returns a value  
    if end-state(s):  
        return utility(s)  
    v = -INF  
    for s in successors(s):  
        v = max(v, min-val(s, alpha, beta) )  
        if v >= beta:  
            return v  
        alpha = max(alpha, v)  
    return v
```



# Alpha-beta pseudocode, pt 2

```
def min-val(s, alpha, beta): #returns a value
    if end-state(s):
        return utility(s)
    v = +INF
    for s in successors(s):
        v = min(v, max-val(s, alpha, beta) )
        if v <= alpha:
            return v
        beta = min(beta, v)
    return v
```