



Artificial Intelligence Programming

Reinforcement Learning

Cindi Thompson

Department of Computer Science
University of San Francisco

Previously ...

- Tuesday we saw how to solve a Markov Decision Process.
- This gives us an optimal policy
 - Mapping of states to actions
 - Lets our agent act optimally (in expectation) in stochastic environments.
- Value Iteration and Policy Iteration assume a lot of available knowledge.
 - All rewards, all state transitions
- Can we learn the values of states without this information?


Reinforcement Learning

- So far, most of the learning algorithms we've looked at have been *supervised*, *passive*, and *offline*.
 - We're given a labeled dataset and need to construct a hypothesis.
- Sometimes our agent is not that lucky.
 - Placed in an environment and forced to learn the best action to take in each state.
 - The environment provides a reward to the agent, but never tells it the best thing to do.
 - The agent must select actions to take.

Reinforcement Learning

- This approach to learning is called *reinforcement learning*.
- Not too different from the way children or animals are trained.
 - Reward positive actions
 - Punish negative actions

Reinforcement Learning Scenario



RL Scenario

- Don't know anything! Unknown environment, transition model, and reward function
- Need to explore the world. General approach: at each state the agent:
 - Selects an available action (it at least knows that!)
 - Receives reward (it can still sense rewards)
 - Observes resulting state (can still sense current state)
 - Repeat until a terminal state

Goal: find an optimal policy from these observations

Reinforcement Learning

Let's start simpler

- How could we do this in a deterministic, episodic environment?
 - Each action leads to a single reward.

Reinforcement Learning

- How could we do this in a deterministic, episodic environment?
 - Each action leads to a single reward.
 - Try each action once, then always select the action with the highest reward.

Reinforcement Learning

- How could we do this in a deterministic, sequential environment?
- Case 1:
 - A sequence of actions leads to a reward at the end.
 - Try each sequence once, then always select the sequence with the highest reward.
- Case 2:
 - Each action has an immediate reward.
 - Use search to find the optimal sequence of actions.

Reinforcement Learning

- What about in a stochastic, episodic environment?
 - Each action has multiple possible outcomes.

Reinforcement Learning

- What about in a stochastic, episodic environment?
 - Each action has multiple possible outcomes.
 - We try each action multiple times and learn the expected utility of each action.
 - This is sometimes called a *bandit problem*

Reinforcement Learning

- What about in a stochastic, episodic environment?
 - Each action has multiple possible outcomes.
 - We try each action multiple times and learn the expected utility of each action.
 - How many times is “enough”?
 - Often, our agent is doing *active learning*; must integrate learning and performance.

Exploration

- Issue: The agent would always like to choose the action with the highest reward.
- This means taking the action with the highest expected utility.
- But if the agent never tries 'bad-looking' actions, it might be missing out on a higher reward.
 - Example: consider two levers:
 - one pays \$1 every time, and
 - the other pays \$0 90% of the time, and \$100 10% of the time.

Exploration

- Issue: early on, we would like to try lots of “non-optimal” actions, as our estimates are probably very inaccurate.
 - This process is called *exploration*.
- As our estimates of the reward for each action get more accurate, we want to select the best action more frequently.
 - This process is called *exploitation*.
- How to do this in a principled way?

Boltzmann exploration

- One way to do this is using Boltzmann exploration.
- Let $U'(a)$ be our estimated utility for taking action a .
- We take an action with probability:

$$\frac{\exp(\frac{U(a)}{k})}{\sum_a \exp(\frac{U(a)}{k})}$$

- Where k is a temperature parameter.
 - k starts high and gradually decreases.
- How will this behave with $k = 1$? $k \ll 1$?

Sequential Problems

- This works great for episodic problems, but what about sequential problems?
- Now we need to think not only about the immediate reward an action generates, but also the value of subsequent states.
- How did we solve this with full information?

Markov Decision Process

- With full information, we modeled this problem as a Markov Decision Process.
- We knew all the states, the transition probabilities between each state, and the rewards for each state.
- We then used value iteration to estimate utilities for each state, or policy iteration to find the policy directly.

Example Utilities

3	0.812	0.868	0.918	+ 1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388
	1	2	3	4

- Value iteration can find utilities for each state, and we use these to construct a policy.

Example Policies

1 0.23 →	2 0.45 →	3 0.68 →	+1
4 0.06 ↑		5 0.33 ↑	-1
6 -0.03 ↑	7 -0.01 →	8 0.13 ↑	9 -0.07 ←

- Policy iteration lets us find policies directly.

Reinforcement learning

- Both value iteration and policy iteration assume a great deal of knowledge about the world.
 - Reward function
 - Transition function
- What if we don't have a model?
- All we know is that there are a set of states, and a set of actions.

Model Estimation

We don't even know the transition function! $P(s'|s,a)$. Might approximate by:

$$\frac{\# \text{transitions } s \rightarrow s' \text{ for } a}{\# \text{times}}$$

- For example, we tried “up” from (1,1) 10 times, and we observe:
 - 8 times we get to (1,2) and $R = -.04$
 - 2 times we get to (2,1) and $R = -.04$
- Now $P((1,2) | (1,1), \text{“up”}) = 8/10$
- and $R(s)$ for both states = $-.04$

Continued exploration will give increasingly accurate estimates. Do value / policy iteration as before.

Model-free learning

- This is pretty slow!
- We really want to learn an optimal *policy*, don't really care about utilities.
- This is called model-free learning: don't even have to learn transition function

Q-learning

- Learning a policy directly is difficult.
- Problem: our data is not of the form: $\langle \text{state}, \text{action} \rangle$
- Instead, it's of the form s_1, s_2, s_3, \dots, R .
- Since we don't know the transition function, it's also hard to learn the utility of a state.
- Instead, we'll learn a function $Q(s, a)$. This will estimate the “utility” of taking action a in state s .

Q-learning

- More precisely, $Q(s, a)$ will represent the value of taking a in state s , then acting optimally after that.
- $Q(s, a) = R(s, a) + \gamma \max_a \sum_{s'} P(s'|s, a) U(s')$
 - γ is our discount factor
- The optimal policy is then to take the action with the highest Q value in each state.
- Of course, we don't know what transition probabilities or utilities are.

Learning the Q function

- To learn Q , we need to be able to estimate the value of taking an action in a state even though our rewards are spread out over time.
- We can do this iteratively.
- Notice that $U(s) = \max_a Q(s, a)$
- We can then rewrite our equation for Q as:
- $Q(s, a) = R(s, a) + \gamma \max_{a'} Q(s', a')$

Learning the Q function

- Let's denote our estimate of $Q(s, a)$ as $\hat{Q}(s, a)$
- We'll keep a table listing each state-action pair and estimated Q -value
- the agent observes its state s , chooses an action a , then observes the reward $r = R(s, a)$ that it receives and the new state s' .
- In the simplest case, it then updates the Q -table according to the following formula:

$$\hat{Q}(s, a) = r + \gamma \max_{a'} \hat{Q}(s', a')$$

Learning the Q function

- The agent uses the estimate of \hat{Q} for s' to estimate \hat{Q} for s .
- Notice that the agent doesn't need any knowledge of R or the transition function to execute this.
- But why put so much reliance on what just happened?
- so we introduce a *learning rate* α that trades off between what we have already learned and the action just taken. Lower over time:

$$\hat{Q}(s, a) = (1 - \alpha) \times \hat{Q}(s, a) + \alpha \times (r + \gamma \max_{a'} \hat{Q}(s', a'))$$

Convergence

- Q-learning is guaranteed to converge as long as:
 - Rewards are bounded
 - The agent selects state-action pairs in such a way that it sees each infinitely often.
 - This means that an agent must have a nonzero probability of selecting each a in each s as the sequence of state-action pairs approaches infinity.

Exploration

- Of course, these are theoretical guarantees.
- In practice, we must sample each state “enough”.
- We can do this with Boltzmann exploration.

$$\frac{\exp(\frac{\hat{Q}(s,a)}{k})}{\sum_a \exp(\frac{\hat{Q}(s,a)}{k})}$$

- In the next slide, we use “bestAction(s2)” as an abbreviation of $\max_{a'} \hat{Q}(s', a')$

Pseudocode

```
Initialize Q table randomly
s = randomState
while not done :
    a = selectAction(s) # use Boltzmann here
    s2 = takeAction(s,a)
    r = reward(s2)
     $Q(s,a) = (1-\alpha) * Q(s,a) + \alpha * (reward + \gamma * \text{bestAction}(s2))$ 
    if s2 is goal
        s = randomState
    else
        s = s2
```

Pros and cons

- Q-learning has proved to be very useful in some settings.
 - No knowledge of problem dynamics needed.
 - No labeled training data needed.
 - Agent can integrate learning and execution.
- It also has some weaknesses
 - Convergence can be slow
 - Q-table is very large
 - Not much generalization

Incorporating generalization

- A big weakness of Q-learning as compared to other learning algorithms is that knowing what to do in state s tells me nothing about how to act in states that are very similar to s .
 - Q-learning has poor *generalization*
 - This is a result of storing all the state-action pairs as a table.
- A standard way to deal with this is to store the table in a *function approximator*.
 - A construct that maps $\langle state, action \rangle$ to reward.

Credit Assignment

- Q-learning works best in environments in which reward is immediate.
- In delayed-reward environments, the *credit assignment problem* becomes an issue.
- To address this, a more sophisticated method known as *TD-learning* can be used.
 - Idea - look ahead to the best action in the next state and use its EU.

Summary

- Policies tell our agent how to act optimally in a stochastic world.
- Reinforcement learning lets us learn a Q function, which gives us a policy.
- No knowledge of the world needed - we just have to be able to select actions and observe rewards.
- This is active, online learning.
- Most extreme example of trading knowledge for time to learn.